

PoC Closure Report

BK-Ingest — AI-Powered Catalog Onboarding · Final engagement document

Prepared by	ADLM Studio — Adedolapo Quasim (engineer of record & ratification authority)	Date	18 July 2026
Client	BuildersKonnnect (BK)	Engagement window	Brief + materials 17 Jul · build complete 18 Jul · demo 31 Jul 2026
Repository / Live	github.com/adedola12/Bk-PoC · bk.adlmstudio.com · api-bk.adlmstudio.com	Verification	Tested end-to-end on the live deployment by the engineer of record, 18 Jul 2026
Reports against	Acceptance Tests & Evaluation Strategy v1.0 (ratified 17-Jul-2026) — the version in force; no threshold amendments were made at any point	Companions	Final Build Report & API Cost Analysis v1.0 · KB v1.1 · Design Doc v1.0 · Addendum D8–D11

1. Closure Statement

The BuildersKonnnect AI Catalog Onboarding proof of concept is **complete, live, and verified against the ratified acceptance criteria**. The engagement asked whether AI can accurately extract, classify, map and transform real manufacturer catalogs into BK’s bulk-upload template with minimal human intervention. The answer, demonstrated on BK’s own sample set and reproducible from the repository, is yes — with the two integrity guarantees intact throughout: nothing wrong entered the catalog silently, and nothing was dismissed silently. The system was delivered thirteen days inside the committed two-week window, every threshold was met without a single ratified amendment, and the evaluation cannot be accused of retro-fitting: the tests were fixed on 17 July, before a line of pipeline code existed.

2. Final Scorecard vs. Ratified Thresholds (T1–T8)

Test	Ratified threshold (v1.0, 17-Jul)	Final result	Verdict
T1 Document triage	10/10 sample files correctly routed	10/10	PASS
T2 Field-level accuracy	≥95% clean · ≥85% hostile	100% clean sources after Stage-5 brand resolution (94.7% interim honestly retained in record); hostile within threshold on sampled sets	PASS
T3 Zero-touch rate	≥70% overall, per catalog	100% (Twyford · Alca · Jaquar · Artize all 100)	PASS
T4 Safety — no silent errors	0 leaks	Real anomalies A1 (“2700 V”) and A3 (part no. on 3 products) caught from source · 10/10 seeded errors flagged · QRG text path rejected by pairing QA	PASS
T5 Pricing policy	100%; history preserved, no overwrites	Append-only enforced at model layer; history queryable; cross-brand comparison shown	PASS
T6 Variant policy	100% of declared options	Base + 9 finishes per product; all derived SKUs labelled derived_unverified	PASS
T7 Throughput	104-pp handbook < 1 hr	2.0 minutes (89 products, 52 calls) — 30× inside threshold	PASS
T8 No silent dismissal	100%; SG_BK.docx designated case	100% — D11 prompt on every non-product route; SG_BK parsed to Satkay vendor profile	PASS

Result: **8 of 8 ratified tests passed.** Verification basis: engineer-of-record end-to-end testing on the live deployment (18-Jul) on top of the automated harness (57/57 tests; scorecards in runs/). Per the evaluation-day runbook, the final formal artifact is one clean `npm run eval -- --full` on 31 July with ground-truth statuses at `human_verified` — a re-confirmation of the results above, generated fresh in front of the client if desired. Supporting: 42/42 rows emitted at 100% zero-touch · 21 descriptions, 0 fact-check failures · Twyford golden round-trip lossless (21×30, zero diffs).

3. Master Checklist — Closed

The ratified strategy (§5) fixed this checklist as the single basis of evaluation. All 25 lines are closed; evidence noted inline.

A. Intake & Triage (T1, T8)	Status
[X] All 10 sample files uploaded through the system UI and classified — <i>10/10 routing</i>	DONE
[X] SG_BK.docx → vendor_document; verification prompt shown; parsed into Satkay vendor profile — <i>RC 196060 · 10-yr warranty · MOQ 1 · 48-hr delivery captured</i>	DONE
[X] Both JPGs → product_image and correctly paired — <i>toilet roll holder · towel rail</i>	DONE
[X] Zero files finalized as non-product without the D11 verification prompt — <i>enforced in Stage 0</i>	DONE
B. Extraction & Normalization (T2)	Status
[X] Alca AM101: decimal-comma normalization, range preserved, pack qty, EAN, WRAS cert, EN norms, 2/15-yr warranty — <i>13,45 kg → 13.45</i>	DONE
[X] Jaquar & Artize: codes, materials, finish lists, drawing dimensions captured and labelled derived	DONE
[X] Bosch HD Handbook segmented at scale with part-number dedupe — <i>89 products, 52 calls</i>	DONE
[X] Bosch QRG vision path triggered by pairing-QA failure — <i>both pages forced to vision</i>	DONE
[X] Every field carries value, unit, raw, source ref, confidence, method — <i>IPR contract</i>	DONE
C. Classification, Brand & IDs (D1, D2, D9)	Status
[X] All products at correct taxonomy nodes — <i>3-method classifier, 3/3</i>	DONE
[X] ADLM cat_ ID register with provenance labels adlm_registry vs bk_confirmed — <i>crosswalk to 4 confirmed BK IDs</i>	DONE
[X] Brand registry: Jaguar→Jaquar alias · Artize ⊂ Jaquar via QUA- signature · Twyford, Alcadrain, Bosch with evidence	DONE
[X] No taxonomy ID guessed; unmapped BK IDs flagged as todos — <i>A7 protection held</i>	DONE
D. Variants, Pricing & Ledger (T5, T6; D4, D5, D10)	Status
[X] Jaquar + Artize expand base + 9 finish variants each, derived-unverified labelled — <i>20 variant rows</i>	DONE
[X] All sample rows blank price + todo entries — <i>no prices exist in sources</i>	DONE
[X] Simulated re-upload appends price events only; history queryable — <i>model-layer overwrite rejection proven</i>	DONE
[X] Price comparison across ≥2 brands within one category — <i>demonstrated in Price Compare</i>	DONE
E. Content Generation & Emission (D3, D6)	Status
[X] Descriptions ~100–130 words + 8–9 tags in golden house style for every emitted row — <i>21 generated</i>	DONE
[X] Fact-check: zero invented facts — <i>0 failures · 1 auto-regeneration · checker rejects poisoned copy (self-test)</i>	DONE
[X] Exact BK template structure; Twyford round-trip lossless — <i>21×30, zero diffs</i>	DONE
[X] 500-row sharding; Cloudinary URLs in media columns	DONE

F. Safety & Reporting (T3, T4, T7)	Status
[X] All known + seeded anomalies flagged; zero silent leaks — <i>A1, A3 from real source</i>	DONE
[X] Review queue shows source region side-by-side per flag	DONE
[X] Final report with per-catalog accuracy, zero-touch, timing + T1–T8 scorecard — <i>Run Report page + this document</i>	DONE
[X] Evaluation harness re-runnable end-to-end in repo — <i>npm run eval → runs/</i>	DONE

4. Deliverables Register

Deliverable	Location / form
Working system — pipeline (Stages 0–10), 7-page showcase UI, deployed live	GitHub main · Vercel client · Render API
Governance corpus — Knowledge Base v1.0/v1.1, Design Document v1.0, Decisions Addendum D8–D11, DECISIONS.md (D12–D19; D20 commissioned)	repo docs/ · Dropbox /BuildersKonnnect (v1.0-era copies) · local files
Acceptance Tests & Evaluation Strategy v1.0 (ratified) + this Closure Report	repo docs/ · PDF
Progress & Achievement Summary · Final Build Report & API Cost Analysis (incl. §8 cost model and funding plan)	PDF
Evaluation harness & regression suite (T1–T8 scorers, anomaly injector, ground-truth schema)	server/eval · npm run eval
ADLM category ID register + BK crosswalk · brand registry · trade taxonomy — adoptable by BK wholesale	MongoDB collections + seed scripts
Easter egg (D20): “Konnnect Concierge” catalog agent — chat search, grounded product Q&A, image search, purchase-intent enquiries, support handoff	Addition prompt docs/CONCIERGE_PROMPT.md · additive Milestone F, gated by its own T9/T10 scorers outside the ratified block

5. Decision Trail — Final State

Twenty logged decisions with unbroken numbering from project chat into the repository: D1–D7 (method: taxonomy, brand registry, generated content, pricing, variants, template-as-contract, triage-first) · D8–D11 (platform: MERN/Anthropic/Cloudinary, ID register, append-only price ledger, no silent dismissal) · D12–D16 (build: triage labels, repo layout, classifier precedence, semantic tier, weblink ingestion) · D17/D18 (proposed, non-blocking) · D19 (open backlog: re-upload SKU reconciliation) · D20 (Konnnect Concierge, additive easter egg). Ratification authority rested with Adedolapo Quasim throughout; no decision was applied unratiified except where explicitly marked proposed.

6. Open Items at Closure (transparent handover)

#	Item	Disposition
1	Evaluation-day formal run (31 Jul): credits top-up (\$25–50), ground-truth statuses to human_verified, npm run eval --full, Render warm-up	Runbook in Final Build Report §9
2	D19 re-upload SKU reconciliation; headless-browser weblink fetcher; per-category BK templates (client dependency); review-UI hardening	Commercial-phase backlog
3	D17/D18 formal ratification; D20 Concierge completion per its own F1–F3 gates and T9/T10 scorers	Non-blocking; additive
4	Production API funding at scale: \$550–950/month optimized (Batch + caching + Haiku tier) at ~50,000 products/month; per-run cost logs as the governance mechanism	Cost model in Final Build Report §8

7. Recommendation & Commercial Path

ADLM’s recommendation to BuildersKonnnect: adopt the PoC’s architecture as the production foundation. The registries compound (every confirmed brand, alias, mapping and review resolution reduces future review volume), the evaluation harness converts to a permanent regression suite, and the decision log gives BK a fully auditable account of why the system behaves as it does. The commercial phase — BK import-API

integration, full supplier-base onboarding, and the \$6 backlog — proceeds on a passed PoC per the original engagement framing. ADLM Studio stands ready to continue as engineering partner.

8. Sign-off

Role	Name	Signature	Date
Engineer of record / ratification authority (ADLM)	Adedolapo Quasim		
Accepted for BuildersKonnnect			

This report closes the proof-of-concept engagement as scoped. Results are stated against Acceptance Tests v1.0 as ratified 17-Jul-2026, unamended. Every metric is regenerable from the repository: npm test · npm run eval · npm run probe. Prepared 18-Jul-2026 by ADLM Studio with Claude (Anthropic) as build assistant.