

Acceptance Tests & Evaluation Strategy

AI-Powered Catalog Onboarding — Proof of Concept

Prepared by	ADLM Studio — Adedolapo Quasim	Date	17 July 2026
Client	BuildersKonnnect (BK)	Version	1.0
Ratification authority	Adedolapo Quasim (ADLM) — interim, effective during build, unless otherwise stated. BK may co-ratify or amend at any point before evaluation day.	References	KB v1.1; Design Doc v1.0; Addendum D8–D11

1. Ratification Message (for the BK team record)

Following your brief and the materials received on 17 July 2026, your stated success criterion — “build a system where we upload files, check if it’s a catalogue or product details, and files it into inventory” — has been formalized into eight measurable acceptance tests (T1–T8, Section 3). Your own brief asks for an evaluation that is measurable rather than impressionistic; these tests are how we make that binding on ourselves. The thresholds below are hereby ratified as the working pass/fail basis for the PoC. They are ratified on an interim basis by ADLM to keep the two-week build moving; BK is invited to confirm, tighten, or amend any threshold at any time before evaluation day, and any amendment will be versioned in this document. What will not change after evaluation begins is the test set itself — results will be reported against the tests as ratified, including any failures, so the evaluation cannot be retro-fitted to the outcome.

Evaluation will run on the five representative catalogs and supporting files you provided, against human-verified ground truth prepared in advance (Section 4). The full metrics report — field-level accuracy, zero-touch rate, safety catches, and processing time per catalog — will accompany the working demonstration.

2. Scope of Evaluation

In scope: document triage of all 10 sample files; extraction, normalization, classification, brand resolution, variant expansion, attribute mapping, description/tag generation, validation gating, and emission to the BK bulk upload template; vendor-document parsing with snapshot verification (D11); pricing todo-flagging and price-event ledger behavior (D4/D10). Out of scope for the PoC: integration with the BK codebase or import API, production review UI polish, and onboarding of suppliers beyond the provided sample set — per the client brief, these follow a passed PoC.

3. Ratified Acceptance Tests (T1–T8)

ID	Test	Metric	Ratified threshold
T1	Document triage	Correct routing of each of the 10 sample files (catalogue / datasheet / image / vendor document / other)	10/10 correct
T2	Field-level accuracy	Required template fields correct vs. human-verified ground truth, per catalog	≥ 95% clean sources (Alca, Jaquar, Artize datasheets; Twyford check) · ≥ 85% hostile sources (Bosch QRG; HD Handbook sampled)
T3	Zero-touch rate	% of product rows emitted with no human intervention (all required fields pass the confidence gate)	≥ 70% overall; reported per catalog
T4	Safety — no silent errors	Known-anomaly classes (A1–A3: wrong units, scrambled label-value pairs, duplicate part numbers) leaking into emitted output unflagged	0 leaks — every seeded/known anomaly must be caught and routed to review
T5	Pricing policy (D4/D10)	Missing prices → blank + todo entry; simulated re-upload with prices → new price event appended, all other fields untouched	100% of applicable rows; price history preserved (no overwrites)
T6	Variant policy (D5)	Declared purchasable options expand to one row per variant, derived SKUs labelled derived-unverified	100% of declared options (Jaquar/Artize: base + 9 finishes each)
T7	Throughput	End-to-end processing time per catalog on PoC hardware	Reported per catalog · single datasheets in minutes · 104-pp handbook target < 1 hr
T8	No silent dismissal (D11)	Every non-product classification surfaces the snapshot-verification prompt (document preview + extracted summary) before filing	100% — SG_BK.docx is the designated test case

Threshold rationale: 95/85 field-accuracy split reflects verified source difficulty (the Bosch QRG demonstrably scrambles label–value pairs under linear extraction); 70% zero-touch matches the commercial goal of minimal human intervention while keeping the review queue as the designed, honest destination for hostile sources; T4 and T8 are absolute because they are integrity tests, not accuracy tests — a single silent error or silent dismissal defeats the system's purpose.

4. Evaluation Strategy — How We Achieve and Prove It

4.1 Ground truth first, build second

Before the pipeline processes anything, a human-verified ground-truth dataset is annotated from the five catalogs: every product, every required template field, correct taxonomy node, correct brand (incl. Artize ⊂ Jaquar lineage), declared variants, and the location of each fact in the source (file/page/region). The golden Twyford file serves as ground truth for the emitter. Annotating first prevents the most common PoC failure: grading the system against answers derived from its own output.

4.2 Staged evaluation gates

Gate	When	What must be true to proceed
G1 — Triage	End of build days 1–2	T1 and T8 pass on all 10 files (incl. SG_BK.docx verification prompt)
G2 — Extraction quality	End of week 1	T2 met on the three clean datasheets; Twyford round-trip check passes; interim accuracy report issued
G3 — Full pipeline	Mid week 2	T5, T6 pass; classification + brand registry + ID register operating; QRG anomalies caught (T4)

Gate	When	What must be true to proceed
G4 — Evaluation day	End of week 2 (target 31 Jul 2026)	All T1–T8 measured and reported, incl. per-catalog T3 and T7; demonstration delivered

Gates fail forward transparently: a missed gate produces a written variance note (what missed, why, corrective action) rather than a moved threshold. Thresholds only change by ratified amendment, never retrospectively.

4.3 Measurement mechanics

Field accuracy = correct required fields ÷ total required fields per catalog, scored automatically against the annotation set with human adjudication of disputes. Zero-touch = rows auto-emitted PASS ÷ total ground-truth rows. Safety = seeded-anomaly recall (the known A1–A3 defects plus 5 synthetic injected errors per hostile catalog — wrong units, swapped values, duplicate codes — all of which must be flagged). Throughput = wall-clock per catalog, recorded by the pipeline itself. Every number in the final report is reproducible by re-running the evaluation harness, which ships inside the PoC repo as a regression suite.

5. Evaluation Basis — Master Checklist

The demonstration is judged complete when every line below is checked. This checklist is the single basis of evaluation; nothing off-list is graded, nothing on-list is skipped.

A. Intake & Triage (T1, T8)	Done
☐ All 10 sample files uploaded through the system UI and classified	
☐ SG_BK.docx → vendor_document; snapshot-verification prompt shown; parsed into Satkay vendor profile (company, RC, contacts, warranty 10yr, MOQ 1, 48hr delivery)	
☐ Both JPGs → product_image and correctly paired (toilet roll holder; towel rail)	
☐ Zero files finalized as non-product/irrelevant without the D11 verification prompt	
B. Extraction & Normalization (T2)	Done
☐ Alca AM101: single product; decimal-comma values normalized (13,45 kg → 13.45 kg); pressure range preserved as range; pack qty parsed (1 20 pcs); EAN, WRAS cert no., EN norms, 2/15-yr warranty captured	
☐ Jaquar & Artize datasheets: product codes, descriptions, materials, finish list, drawing dimensions captured; drawing-derived dimensions labelled derived	
☐ Bosch HD Handbook: products segmented at scale; part-number dedupe applied	
☐ Bosch QRG: vision-mode extraction path triggered by pairing-QA failure	
☐ Every extracted field carries value, unit, source reference, confidence, extraction method	
C. Classification, Brand & IDs (D1, D2, D9)	Done
☐ All products classified to correct trade taxonomy nodes (Twyford→Sanitary Ware; Alca→Concealed Cisterns; Jaquar/Artize→Bathroom Accessories; Bosch→Tools & Equipment sub-trades)	
☐ ADLM category ID register issues cat_-format IDs for every node; ID provenance marked (adlm_registry vs bk_confirmed)	
☐ Brand registry resolves: Jaguar→Jaquar alias; Artize ⊂ Jaquar lineage; Twyford, Alcadrain, Bosch registered with evidence	
☐ No taxonomy ID guessed; unmapped BK IDs flagged, not fabricated	
D. Variants, Pricing & Ledger (T5, T6; D4, D5, D10)	Done
☐ Jaquar + Artize expand to one row per finish variant (base + 9 each); derived SKUs labelled derived-unverified	
☐ All sample rows have blank price + todo entries (no prices exist in sources)	
☐ Simulated re-upload with prices: price events appended to ledger; no other field modified; history queryable	
☐ Price-comparison view demonstrable across ≥ 2 brands within one category	

E. Content Generation & Emission (D3, D6)	Done
[] Descriptions (~100–130 words) and 8–9 tags generated in golden-example house style for every emitted row	
[] Post-generation fact-check: every claim in generated copy traces to an extracted field; zero invented facts	
[] Output emitted in exact BK template structure; Twyford golden file round-trips through the emitter without loss	
[] Files shard automatically at the 500-row template cap; Cloudinary URLs populate media columns	
F. Safety & Reporting (T3, T4, T7)	Done
[] All known anomalies A1–A3 and all seeded synthetic errors flagged to review — zero silent leaks	
[] Review queue shows side-by-side source region for every flagged field	
[] Final report: per-catalog field accuracy, zero-touch rate, processing time; overall T1–T8 scorecard	
[] Evaluation harness re-runnable end-to-end inside the repo (regression suite)	

6. Sign-off

Role	Name	Signature	Date
Ratified by (ADLM — interim authority)	Adedolapo Quasim		
Acknowledged / co-ratified by (BK — optional pre-evaluation)			

Amendments to thresholds are made only by a versioned revision of this document prior to evaluation day. Results are reported against the ratified version in force at evaluation.